

innovate

achieve

lead



# Conversational Systems Design

## Lecture 2

Kedar Kanhere  
02/06/2024

**BITS Pilani**  
Pilani Campus

# Session Content but before that



## Quiz and Assignment Tentative Schedule

Exam Type	Start Date	End Date	Marks
Quiz 1	13/07/2024 7:00 PM	14/07/2024 7:00 PM	5 (Best 2 of 3)
Quiz 2	31/08/2024 7:00 PM	01/09/2024 7:00 PM	5 (Best 2 of 3)
Quiz 3	07/09/2024 7:00 PM	08/09/2024 7:00 PM	5 (Best 2 of 3)
Assignment 1	20/06/2024	18/07/2024	15
Assignment 2	25/08/2024	22/09/2024	15

# Session Content but before that

---



Webinar will tentatively be on these dates:

19-Jun-24
10-Jul-24
14-Aug-24
4-Sep-24

# Session Content

---



- Importance of text pre-processing
- Overview of past learnings (NLP)
  - Stemming
  - Tokenization
  - POS Tagging
- Text to Speech
  - Understanding T2S algorithms basic layer
  - Different algorithms working
  - Code and example via different libraries
  - Importance of T2S Synthesis, ASR etc.

# Importance of text pre-processing



- **Text pre-processing** is a crucial step in Natural Language Processing (NLP) that involves cleaning and preparing text data for analysis.
- It includes techniques such as:
  - Tokenization
  - Stop Word Removal
  - Stemming
  - Lemmatization
  - Text Cleaning

Importance: These steps help in converting raw text into a more manageable and analyzable format, facilitating better model performance and more accurate results.

# Tokenization



- Tokenization is the process of splitting text into smaller units called tokens, which can be words, phrases, or symbols.
  - Example: Input: "Natural Language Processing is fascinating."
  - "Output: ["Natural", "Language", "Processing", "is", "fascinating"]"
- Tools
  - NLTK: A leading platform for building Python programs to work with human language data.
  - SpaCy: An open-source software library for advanced NLP in Python.
  - Tokenizer APIs: Available in various programming languages and platforms.
- Importance: Tokenization is the first step in the text pre-processing pipeline and is crucial for the performance of subsequent steps.

# Tokenization



**Inconsistent Data Representation** Without tokenization, text data remains as large, unmanageable strings, making it difficult to analyze and process.

**Model Performance:** Machine learning models require numerical or categorical input. Without tokenization, converting text into a suitable format is impossible, leading to poor model performance.

**Difficulty in Feature Extraction:** Tokenization allows for extracting features such as word frequencies, n-grams, and more. Skipping this step hinders effective feature extraction.

**Ineffective Text Cleaning:** Tokenization is often the first step in text cleaning. Without it, removing stop words, punctuation, and performing stemming/lemmatization becomes challenging.

**Error Propagation:** Errors in initial steps propagate through the pipeline, leading to inaccuracies in tasks like sentiment analysis, NER, and POS tagging.

# Stop-Words



- Stop words are common words that are usually ignored in text processing because they do not carry significant meaning.
  - Examples: Common stop words include "a", "the", "and", "in", "to", etc.
- Purpose:
  - Noise Reduction: Removing stop words helps in reducing the noise in the text data.
  - Efficiency: It reduces the size of the text data, making processing faster and more efficient.
- Tools:
  - NLTK: Provides a predefined list of stop words and functions for their removal.
  - SpaCy: Offers built-in support for stop word removal in various languages.
- Example:
  - Input: "The quick brown fox jumps over the lazy dog."
  - Output: "quick brown fox jumps lazy dog"
- Importance: Removing stop words helps in focusing on the words that are more likely to be significant in the analysis, thereby improving the performance of NLP models.



# Stemming and Lemmatization



- Stemming: Reduces words to their base or root form by removing suffixes. It may not always produce a real word.
- Lemmatization: Converts words to their base form (lemma) using morphological analysis. It always returns a valid word.
- Examples:
  - Stemming: "running" -> "run", "jumps" -> "jump"
  - Lemmatization: "better" -> "good", "running" -> "run"
- Tools:
  - NLTK:
    - PorterStemmer for stemming
    - WordNetLemmatizer for lemmatization
  - **SpaCy**: Offers built-in lemmatization capabilities.

# Stemming and Lemmatization



- Differences:
  - Accuracy: Lemmatization is generally more accurate than stemming.
  - Complexity: Stemming is simpler and faster but less accurate.
- Use Case: Choose stemming for quick and dirty text processing; use lemmatization for tasks requiring higher accuracy.
- Importance: Both techniques help in normalizing words to their base forms, which reduces the dimensionality of the text data and improves the performance of NLP models.
- Problems if we don't do stemming or lemmatization:
  - High dimensionality
  - Inconsistent Data
  - Reduced model performance and difficulty in text analysis
  - Lower accuracy in search and retrieval

# Case Normalisation



- Case normalization is the process of converting all characters in the text to a uniform case, either lower case or upper case, to ensure consistency.
- Purpose:
  - Consistency: Ensures that words are treated equally regardless of their case.
  - Reduction of Redundancy: Helps in reducing redundancy by treating "Apple" and "apple" as the same word.
- Example:
  - Input: "Natural Language Processing"
  - Output: "natural language processing"
- Tools:
  - Python String Methods: `.lower()` and `.upper()`
  - NLTK: Provides functions for case normalization.
  - SpaCy: Built-in support for case normalization.
- Importance:
  - Improves Text Processing: Case normalization simplifies text processing by reducing the number of unique tokens.
  - Enhances Model Performance: Models become more efficient as they deal with fewer variations of the same word.
- Note: Case normalization is particularly useful when the case of the text does not carry significant meaning for the analysis.

# Text Cleaning



- Text cleaning involves removing unwanted elements from the text to make it suitable for analysis.
- Common Techniques:
  - Removing Punctuation: Eliminating characters such as periods, commas, and exclamation marks.
  - Removing Special Characters: Removing symbols like #, \$, %, etc.
  - Removing Numbers: Excluding digits unless they are relevant to the analysis.
  - Removing HTML Tags: Stripping HTML content from web-scraped text.
  - Handling Contractions: Expanding contractions (e.g., "don't" to "do not").
- Tools:
  - Regular Expressions (Regex): Powerful for pattern matching and substitution.
  - NLTK: Provides functions for various text cleaning tasks.
  - SpaCy: Built-in functions for text cleaning.
- Example:
  - Input: "Hello, world! Visit us at <https://example.com> #NLP"
  - Output: "Hello world Visit us at example com NLP"
- Importance:
  - Enhances Data Quality: Cleaned text is more consistent and easier to analyze.
  - Improves Model Accuracy: Cleaner data leads to better-performing models by reducing noise and irrelevant information.

# Named Entity Recognition (NER)



- Named Entity Recognition (NER) is a subtask of information extraction that identifies and classifies named entities in text into predefined categories such as person names, organizations, locations, dates, etc.
  - **Categories:**
    - **Person:** Names of people (e.g., "John Doe")
    - **Organization:** Names of organizations (e.g., "Google")
    - **Location:** Geographical locations (e.g., "Paris")
    - **Date/Time:** Dates and times (e.g., "January 1, 2020")
    - **Others:** Money, percentages, etc.
  - **Purpose:**
    - **Information Extraction:** Helps in extracting structured information from unstructured text.
    - **Improved Search and Retrieval:** Enhances the performance of search engines and information retrieval systems.
  - **Tools:**
    - **SpaCy:** Provides a pre-trained NER model and tools for custom training.
    - **NLTK:** Offers NER capabilities through its chunking module.
    - **Stanford NER:** A widely used NER tool developed by Stanford University.
- Importance:** NER is essential for understanding the context and extracting relevant information from large volumes of text, making it a key component in various NLP applications.

# Named Entity Recognition (NER)



**Rule-based Approaches:** Use predefined rules and patterns to identify entities.

- **Example:** Regular expressions to identify dates and email addresses.
- **Pros:** Simple to implement and interpret.
- **Cons:** Limited flexibility and scalability.

**Machine Learning Approaches:** Use statistical models trained on labeled data to identify entities.

- Conditional Random Fields (CRF)
- Hidden Markov Models (HMM)
- **Pros:** More flexible and accurate than rule-based approaches.
- **Cons:** Require labeled training data and computational resources.

# Named Entity Recognition (NER)



- **Deep Learning Approaches:** Use neural networks to automatically learn features and patterns from data
  - Recurrent Neural Networks (RNN)
  - Long Short-Term Memory (LSTM)
  - Transformer-based models (e.g., BERT)
- **Pros:** High accuracy and ability to capture complex patterns.
- **Cons:** Require large datasets and significant computational power.
- **Example:**
  - **Input:** "Apple is looking at buying U.K. startup for \$1 billion."
  - **Output:** Entities: [("Apple", "ORG"), ("U.K.", "LOC"), ("1 billion", "MONEY")]

## Importance:

Choosing the right technique depends on the specific requirements of the task, available data, and resources. Machine learning and deep learning approaches are preferred for their accuracy and scalability.

# Part of Speech Tagging (POS)



- Part-of-Speech (POS) tagging is the process of assigning a part of speech to each word in a sentence. The parts of speech include nouns, verbs, adjectives, adverbs, pronouns, conjunctions, prepositions, and interjections
- **Understanding Syntax:** Helps in understanding the grammatical structure of sentences.
- **Disambiguation:** Resolves ambiguities by providing context to words (e.g., "book" as a noun vs. "book" as a verb).

**Example** The quick brown fox jumps over the lazy dog."

- POS Tags:  
The (DT) quick (JJ) brown (JJ) fox (NN) jumps (VBZ) over (IN) the (DT) lazy (JJ) dog (NN)

## Tools:

- **NLTK:** Provides a comprehensive POS tagging module.
- **SpaCy:** Offers efficient and accurate POS tagging capabilities.
- **Stanford POS Tagger:** A robust tool developed by Stanford University.

## Importance:

POS tagging is fundamental for many NLP tasks such as parsing, text-to-speech conversion, and information extraction. It enables a deeper understanding of the syntactic and semantic properties of text.



# Part of Speech Tagging (POS)



- **Rule-based Tagging:** Uses a set of hand-crafted rules to assign POS tags.
  - **Example:** "If a word ends in 'ly', tag it as an adverb (RB)."
- **Statistical Tagging:** Uses probabilistic models based on the likelihood of a word's POS tag given its context.
  - Hidden Markov Models (HMM)
  - Maximum Entropy Models
- **Machine Learning Approaches:** Use supervised learning techniques to predict POS tags.
  - Support Vector Machines (SVM)
  - Conditional Random Fields (CRF)
- **Deep Learning Approaches:** Uses neural networks to automatically learn features from data.
  - Recurrent Neural Networks (RNN)
  - Long Short-Term Memory (LSTM)
  - Transformer-based models (e.g., BERT)

## Example:

- Sentence: "The cat sat on the mat."
- POS Tags: The (DT) cat (NN) sat (VBD) on (IN) the (DT) mat (NN)

# Sentiment Analysis



- Sentiment analysis, also known as opinion mining, is the process of determining the sentiment expressed in a piece of text. It classifies text into positive, negative, or neutral sentiments.

**Understanding Public Opinion:** Helps businesses understand customer opinions and feedback.

**Market Research:** Analyzing trends and opinions about products, services, or events.

**Social Media Monitoring:** Tracking sentiment trends on social media platforms.

## Applications:

**Customer Feedback Analysis:** Understanding customer satisfaction and improving services.

**Brand Monitoring:** Tracking public sentiment towards a brand.

**Political Sentiment:** Analyzing public opinion on political issues or candidates.

**Product Reviews:** Assessing the sentiment in product reviews to gauge consumer reactions.

# Sentiment Analysis



## Techniques:

**Lexicon-based Methods:** Use predefined lists of positive and negative words.

**Machine Learning Approaches:** Train classifiers (e.g., SVM, Naive Bayes) on labeled data.

**Deep Learning Approaches:** Use neural networks (e.g., LSTM, BERT) for higher accuracy.

## Example:

Input: "I love the new design of your website!"

Output: Positive

## Importance:

Sentiment analysis provides valuable insights into the emotions and opinions expressed in text, enabling better decision-making and strategy formulation.

# Text to Speech Synthesis



- Text-to-Speech (TTS) synthesis is the process of converting written text into spoken words using computational methods.
- Purpose:
  - Accessibility: Provides access to information for visually impaired users.
  - User Experience: Enhances user interaction in virtual assistants and conversational agents.
  - Automation: Automates tasks that require reading text, such as news reading or announcements.
- Components:
  - Text Analysis:
    - Tokenization: Breaking text into smaller units such as sentences and words.
    - Linguistic Analysis: Determining the part of speech, phonetic transcription, and prosody (intonation, stress, rhythm).
  - Acoustic Modeling:
    - Phoneme Synthesis: Generating speech sounds based on phonetic transcription.
    - Prosody Generation: Adding natural intonation, stress, and rhythm to synthesized speech.
  - Speech Synthesis:
    - Concatenative Synthesis: Combining pre-recorded speech segments.
    - Formant Synthesis: Using mathematical models to generate speech sounds.
    - Waveform Synthesis: Using neural networks to generate high-quality, natural-sounding speech.

# Text to Speech Synthesis



- Tools:
  - Google Text-to-Speech: A widely used TTS engine with high-quality voices.
  - Amazon Polly: A cloud-based service that converts text into lifelike speech.
  - IBM Watson Text to Speech: Provides a range of customizable voices.
  - Open Source Tools: eSpeak, Festival, and MaryTTS
- Example:
  - Input: "Welcome to the world of Text-to-Speech synthesis."
  - Output: (Spoken audio)
- **Importance:** TTS synthesis plays a crucial role in improving accessibility and enhancing user experiences across various applications and devices.

# Text Generation Techniques



Text generation involves creating coherent and contextually relevant text from a given input or set of rules.

## Key Techniques:

Rule-based Systems : Use predefined rules and templates to generate text.

Example: Fill-in-the-blank templates for automated report generation

Markov Chains : Use probabilistic models based on the likelihood of word sequences.

Example: Generating text by predicting the next word based on the previous one

Recurrent Neural Networks (RNNs): Use neural networks with loops to maintain context over sequences.

Example: Generating poetry or short stories

term Long Short-Term Memory Networks (LSTMs) : A type of RNN designed to better handle long-dependencies.

Example: Generating more coherent paragraphs and articles.

Transformer Models: Use self-attention mechanisms to capture long-range dependencies in text.

Example: GPT-3 generating articles, stories, and dialogue.

# State-based and Rule-based Dialogue Systems



Dialogue systems designed to follow specific states or rules to manage interactions with users.

**State-based Dialogue Systems :** Systems that manage conversations using predefined states and transitions between those states.

**Rule-based Dialogue Systems:** Systems that use predefined rules and templates to generate responses and manage interactions.

Examples : If user says "Hello", respond with "Hi, how can I help you today?"

If user asks for account balance, respond with "Please provide your account number."

**Components:**

**Rule Engine:** Processes user input based on predefined rules.

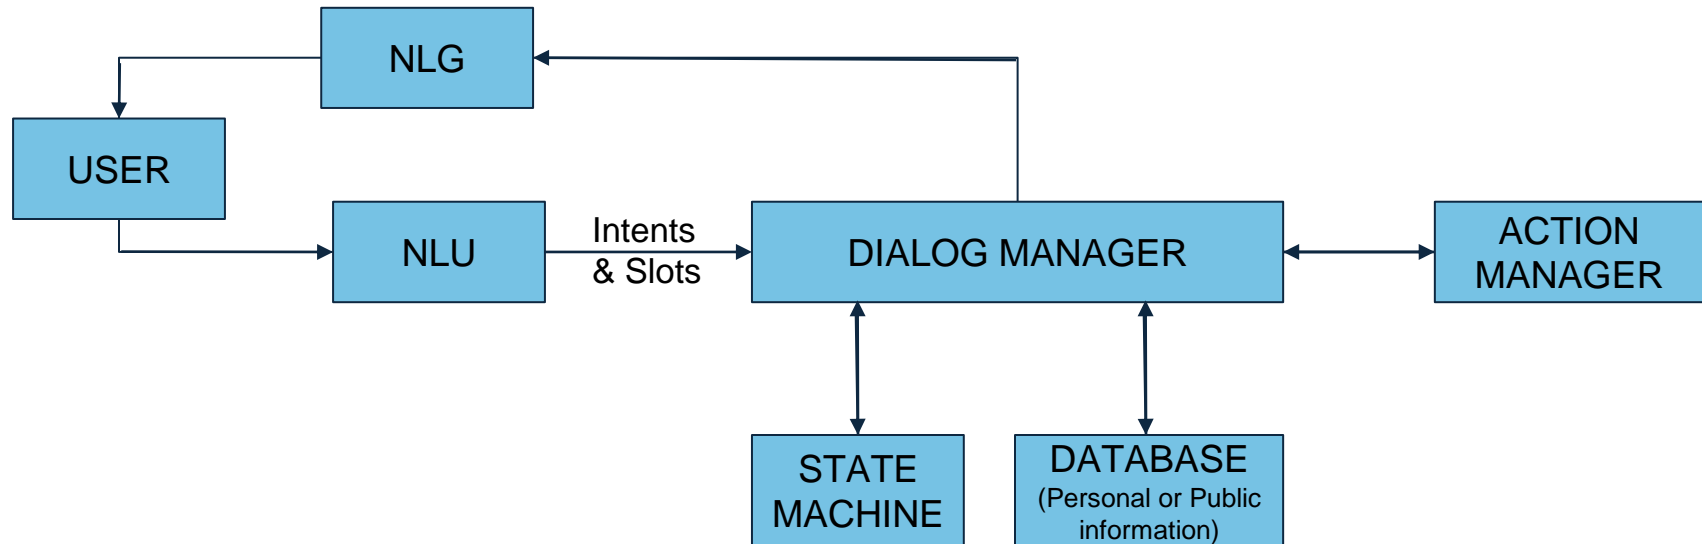
**Template Manager:** Uses templates to generate responses.

**Context Manager:** Maintains context of the conversation to apply relevant rules.

# State-based and Rule-based Dialogue Systems



## State-based Dialogue Systems





# Introduction to Intent Recognition and Slot Filling



Intent Recognition: The process of identifying the goal or purpose behind a user's input in a conversation.

Purpose: Helps in understanding user requests and guiding the conversation appropriately.

Example Intents:

Booking a flight

Slot Filling: Extracting specific pieces of information (slots) from the user's input that are necessary to complete the intent.

Purpose: Provides detailed information required to fulfill the user's request.

Example

Slots for Flight Booking:

Destination

Departure

DateReturn

Date

Number of Passengers

# Introduction to Intent Recognition and Slot Filling



## Components:

Natural Language Understanding (NLU) Module: Processes user input to recognize intents and extract slots.

Dialogue Manager: Uses recognized intents and filled slots to manage the conversation flow and fulfill the user's request.

## Techniques:

Rule-based Methods: Use predefined patterns and templates to recognize intents and extract slots.

Machine Learning Approaches: Train classifiers on labeled datasets to predict intents and extract slots.

Deep Learning Approaches: Use neural networks, particularly sequence-to-sequence models, to handle more complex and varied inputs.

## Example:

User Input: "I want to book a flight to New York on June 5th."

Recognized Intent: Book Flight

Extracted Slots:

Destination: New York

Departure Date: June 5th