

Bias and UI/UX in LLMs

Bias

Bias in LLMs can stem from several sources:

- **Training Data:** LLMs are trained on vast amounts of text data from the internet, which inherently contains societal biases related to race, gender, culture, and religion.
- **Model Specifications:** The architecture and algorithms used can also introduce or amplify biases present in the training data.
- **Algorithmic Constraints:** Certain design choices and constraints in the model's development can lead to biased outcomes.
- **Product Design and Policy Decisions:** Decisions made during the product development phase, including how the model is fine-tuned and deployed, can influence the presence and impact of bias.

Types of Bias

Bias in LLMs can manifest in various forms:

- **Gender Bias:** Associating certain professions or roles with specific genders. For example, tasks related to software development often show a gendered distribution, with technical roles being more frequently associated with men.
- **Cultural Bias:** Models may favor Western cultural norms and fail to appropriately adapt to other cultural contexts. For instance, multilingual models often exhibit biases towards Western culture over others.
- **Religious Bias:** LLMs can display biases against specific religious groups, such as associating Muslims with violence more frequently than other religious groups.
- **LGBTQ+ Bias:** Models may encode biases harmful to the LGBTQ+ community, which can be measured and somewhat mitigated through community-driven benchmarks

Whose Opinions Do Language Models Reflect



language models have
offered subjective opinions
to controversial social and
political queries



whose opinions (if any) do
language models reflect?

How to measure Bias in LLMs?

- There are multiple metrics which can be used to measure bias in LLMs
 - RBS and ABS
 - BiQ
- For fairness the set of metrics used are:
 - Demographic Parity
 - Equality of odds
 - Disparate Treatment
 - Disparate Impact

RBS and ABS

Representative Bias Score (RBS): This metric measures the extent to which LLMs generate outputs that reflect the experiences of certain identity groups over others. It helps in identifying if the model is biased towards specific demographics such as race, gender, or sexual orientation.

Affinity Bias Score (ABS): This metric evaluates the model's preference for specific narratives or viewpoints. It identifies biases in evaluative patterns, often referred to as "bias fingerprints" within the model.

RBS and ABS

- Representative Bias Score (RBS)


Example: Suppose an LLM is used to generate news articles. If the model disproportionately generates articles featuring male scientists over female scientists, despite equal representation in the training data, this indicates a high RBS. The model's outputs reflect the experiences of male scientists more prominently, suggesting a bias towards this demographic.

- Affinity Bias Score (ABS)

Example: Consider an LLM trained to provide movie reviews. If the model consistently rates movies with male protagonists higher than those with female protagonists, this reflects an affinity bias. The ABS would quantify this bias by comparing the average ratings for movies based on the gender of the protagonist.

BiQ

Bias Intelligence Quotient (BiQ): Part of the Comprehensive Bias Neutralization Framework (CBNF), BiQ is a multi-dimensional metric that combines several fairness metrics to assess and mitigate racial, cultural, and gender biases in LLMs. It enhances the Large Language Model Bias Index (LLMBI) with additional fairness metrics, providing a nuanced approach to bias detection and mitigation



Example: An LLM is evaluated using the BiQ framework, which includes metrics like demographic parity and equality of odds. If the model shows that job application recommendations favor male candidates over equally qualified female candidates, the BiQ would highlight this gender bias. The BiQ score would incorporate these fairness metrics to provide a comprehensive assessment of the model's bias.

Group Fairness Metrics

Demographic Parity: This metric ensures that the model's predictions are independent of sensitive attributes like age, gender, or race. It measures whether each demographic group receives similar outcomes.

Equality of Odds: This metric checks if the model's error rates are the same across different demographic groups. It ensures that the model does not disproportionately misclassify certain groups

Group Fairness Metrics

- Demographic Parity

Example: An LLM used in a hiring process should ensure that the proportion of recommended candidates from different demographic groups (e.g., gender, race) is like their proportion in the applicant pool. If the model recommends 60% male candidates when the applicant pool is 50% male, this indicates a violation of demographic parity.

- Equality of Odds

Example: In a medical diagnosis application, an LLM should have similar false positive and false negative rates across different demographic groups. If the model has a higher false negative rate for diagnosing a disease in women compared to men, it violates equality of odds, indicating gender bias in the model's predictions.

Causal Inference and Randomized Experiments

Disparate Treatment and Disparate Impact

- Using causal inference and randomized experiments, these metrics measure whether different groups are treated differently by the model (disparate treatment) and whether the model's decisions disproportionately affect certain groups (disparate impact).
- These methods provide a comprehensive framework for understanding and improving fairness in algorithmic decisions.

Causal Inference and Randomized Experiments

Disparate Treatment

- Example: An LLM used for credit scoring should not treat applicants differently based on their race. If an experiment shows that changing the race of an applicant (while keeping other factors constant) results in a different credit score, this indicates disparate treatment.

Disparate Impact

- Example: An LLM used for loan approvals should not disproportionately reject applications from a specific racial group. If an analysis reveals that applicants from a particular race are rejected at a higher rate than others, this indicates disparate impact, suggesting racial bias in the model.

Mitigating Bias



Bias Identification and Quantification: Developing benchmarks and tools to identify and measure bias in LLMs is a crucial first step.



Data Curation and Augmentation: Curating training data to be more representative and augmenting it with additional data from underrepresented groups can help reduce bias.



Model Fine-Tuning: Fine-tuning models on data written by or about marginalized communities can mitigate specific biases.



Algorithmic Adjustments: Implementing algorithmic techniques such as adversarial training and regularization can help reduce bias.



Community Involvement: Engaging with communities to develop benchmarks and evaluate model outputs can ensure that the models are more equitable and less harmful.

Debiasing

- Counterfactual Data Augmentation (CDA): Adding data that counteracts biases.
- CDA Example: Augmenting a dataset with gender-neutral job descriptions to reduce gender bias in job recommendations.
- Self-Debias: A method that fine-tunes models to reduce bias.
- Self-Debias Example: Fine-tuning a model with gender-balanced text to reduce gender bias.
- Iterative Nullspace Projection: Projects biased components of embeddings to a null space.
- Iterative Nullspace Projection Example: Removing gendered dimensions from word embeddings.

Evaluation Toolkits



FairPy: A toolkit for evaluating and mitigating social biases in LLMs.




Functionality: Provides interfaces to connect bias identification tools with LLMs and test custom models.



Example: Using FairPy to test a custom model for racial bias in sentiment analysis.

Introduction to Conversational UI/UX

- Conversational UX involves designing user experiences that incorporate conversational interfaces such as chatbots and voice assistants.
- It enhances user interaction by making technology more intuitive and responsive.



Principles of Conversational UX

1. Responsiveness: Quick and efficient responses.

2. Availability: 24/7 accessibility across devices.

3. Simplicity: Easy and intuitive interactions.

4. Exit Options: Providing users with a way to exit the conversation or reach human support.

Significance of Conversational UX



USER BENEFITS: IMMEDIATE ISSUE
RESOLUTION, MULTITASKING
CAPABILITIES.



ORGANIZATIONAL BENEFITS: COST
REDUCTION, FREEING HUMAN
RESOURCES FROM MUNDANE TASKS.

Types of Conversational UX

1. Chatbots:
Automated text-
based interactions.

2. Interactive
Applications: Mobile
apps with interactive
features.

3. Digital Voice
Assistants: AI-driven
voice interactions
(e.g., Siri, Alexa).




Example 1: Bank of America's Erica

Erica is a virtual assistant for banking needs, improving customer service and task efficiency.

Example 2: Duolingo

- Duolingo uses interactive exercises for language learning and has over 500 million downloads, making it a popular educational tool.



Example 3: Domino's Facebook Messenger Bot

- This chatbot allows ordering and tracking through Facebook Messenger, streamlining the ordering process.

Conclusion

- Conversational UX is a transformative tool for enhancing user interactions.
- The future holds potential for further advancements and wider adoption.